

Lab Exercise 10- Implementing Resource Quota in Kubernetes

Objective:

In Kubernetes, Resource Quotas are used to control the resource consumption of namespaces. They help in managing and enforcing limits on the usage of resources like CPU, memory, and the number of objects (e.g., Pods, Services) within a namespace. This exercise will guide you through creating and managing Resource Quotas to limit the resources used by applications in a specific namespace.

Step 1: Understand Resource Quotas

Resource Quotas allow you to:

- Limit the amount of CPU and memory a namespace can use.
- Control the number of certain types of resources (e.g., Pods, Services, PersistentVolumeClaims) in a namespace.
- Prevent a namespace from consuming more resources than allocated, ensuring fair usage across multiple teams or applications.

Step 2: Create a Namespace

First, create a namespace where you will apply the Resource Quota. This helps in isolating and controlling resource usage within that specific namespace.

Create a YAML file named ***quota-namespace.yaml*** with the following content:

```
apiVersion: v1
kind: Namespace
metadata:
  name: quota-example # The name of the namespace.
```

```
vim +
3 apiVersion: v1
2 kind: Namespace
1 metadata:
4   name: quota-example    # The name of the namespace.
~
```

Apply the YAML to create the namespace:

kubectl apply -f quota-namespace.yaml

```
Apple ~ .....
> cat << EOF > quota-namespace.yaml
apiVersion: v1
kind: Namespace
metadata:
  name: quota-example
EOF

Apple ~ .....
> kubectl apply -f quota-namespace.yaml
namespace/quota-example created
```

Verify that the namespace is created:

kubectl get namespaces

```
Apple ~ .....
> kubectl get namespaces

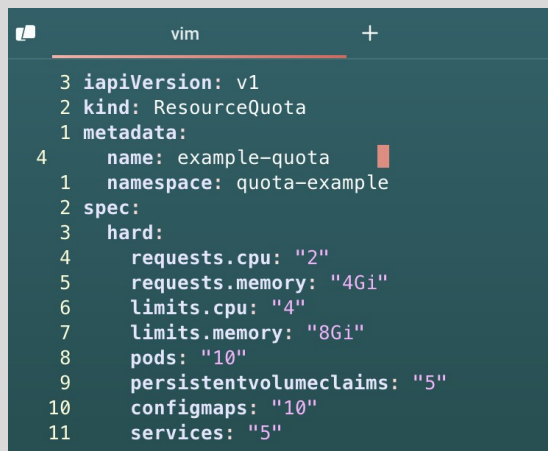
NAME              STATUS   AGE
default           Active   8d
kube-node-lease   Active   8d
kube-public        Active   8d
kube-system        Active   8d
quota-example      Active   28s
```

You should see quota-example listed in the output.

Step 3: Define a Resource Quota

Next, create a Resource Quota YAML file named ***resource-quota.yaml*** with the following content:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: example-quota # The name of the Resource Quota.
  namespace: quota-example # The namespace to which the Resource Quota will apply.
spec:
  hard:
    # The hard limits imposed by this Resource Quota.
    requests.cpu: "2" # The total CPU resource requests allowed in the namespace (2 cores).
    requests.memory: "4Gi" # The total memory resource requests allowed in the namespace (4 GiB).
    limits.cpu: "4" # The total CPU resource limits allowed in the namespace (4 cores).
    limits.memory: "8Gi" # The total memory resource limits allowed in the namespace (8 GiB).
    pods: "10" # The total number of Pods allowed in the namespace.
    persistentvolumeclaims: "5" # The total number of PersistentVolumeClaims allowed in the namespace.
    configmaps: "10" # The total number of ConfigMaps allowed in the namespace.
    services: "5" # The total number of Services allowed in the namespace.
```

A screenshot of a vim editor window. The title bar shows 'vim' and a '+' icon. The editor content is the same YAML file as shown in the previous block, with line numbers 1 through 11 on the left margin. The text is color-coded: 'apiVersion' is blue, 'kind' is blue, 'metadata' is blue, 'name' is blue, 'namespace' is blue, 'spec' is blue, 'hard' is blue, 'requests' is blue, 'cpu' is blue, 'memory' is blue, 'limits' is blue, 'pods' is blue, 'persistentvolumeclaims' is blue, 'configmaps' is blue, and 'services' is blue. The values are in quotes and some are in different colors: '2' is red, '4Gi' is green, '4' is red, '8Gi' is green, '10' is red, '5' is green, '10' is red, and '5' is green.

Step 4: Apply the Resource Quota

Apply the Resource Quota YAML to the namespace:

```
kubectl apply -f resource-quota.yaml
```

```
Apple Home ~ .....  
> kubectl apply -f resource-quota.yaml  
resourcequota/example-quota created
```

Verify that the Resource Quota is applied:

```
kubectl get resourcequota -n quota-example
```

```
Apple Home ~ ..... * docker-desktop © 02:25:04 PM  
> kubectl get resourcequota -n quota-example  
NAME      AGE   REQUEST                                     LIMIT  
example-quota 28s   configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4Gi, services: 0/5 limits.cpu: 0/4, limits.memory: 0/8Gi
```

To see the details of the applied Resource Quota:

```
kubectl describe resourcequota example-quota -n quota-example
```

```
Apple Home ~ .....  
> kubectl describe resourcequota example-quota -n quota-example  
Name:      example-quota  
Namespace: quota-example  
Resource   Used  Hard  
-----  
configmaps 1    10  
limits.cpu  0     4  
limits.memory 0    8Gi  
persistentvolumeclaims 0    5  
pods        0    10  
requests.cpu 0     2  
requests.memory 0    4Gi  
services    0     5
```

Step 5: Test the Resource Quota

Let's create some resources in the quota-example namespace to see how the Resource Quota affects them.

Deploy a ReplicaSet with Resource Requests and Limits

Create a YAML file named ***nginx-replicaset-quota.yaml*** with the following content:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  namespace: quota-example
spec:
  replicas: 5          # Desired number of Pod replicas.
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
      resources:      # Define resource requests and limits.
        requests:
```

```
memory: "100Mi"
cpu: "100m"
limits:
  memory: "200Mi"
  cpu: "200m"
```

```

> cat << EOF > nginx-replicaset-quota.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  namespace: quota-example
spec:
  replicas: 5           # Desired number of Pod replicas.
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
          resources:      # Define resource requests and limits.
            requests:
              memory: "100Mi"
              cpu: "100m"
            limits:
              memory: "200Mi"
              cpu: "200m"
EOF
```

Explanation:

This ReplicaSet requests a total of 500m CPU and 500Mi memory across 5 replicas. It also limits each replica to use a maximum of 200m CPU and 200Mi memory.

Apply this YAML to create the ReplicaSet:

```
kubectl apply -f nginx-replicaset-quota.yaml
```

```

er-desktop 02:27:46 PM
> kubectl apply -f nginx-replicaset-quota.yaml
replicaset.apps/nginx-replicaset created
```

Check the status of the Pods and ensure they are created within the constraints of the Resource Quota:

```
kubectl get pods -n quota-example
```

```

> kubectl get pods -n quota-example

NAME                                READY   STATUS              RESTARTS   AGE
nginx-replicaset-48stw              0/1     ContainerCreating    0           29s
nginx-replicaset-52g9d              1/1     Running              0           30s
nginx-replicaset-bl956              1/1     Running              0           30s
nginx-replicaset-lrnjn              1/1     Running              0           30s
nginx-replicaset-p2w8d              1/1     Running              0           29s
```

To describe the Pods and see their resource allocations:

```
kubectl describe pods -l app=nginx -n quota-example
```

```

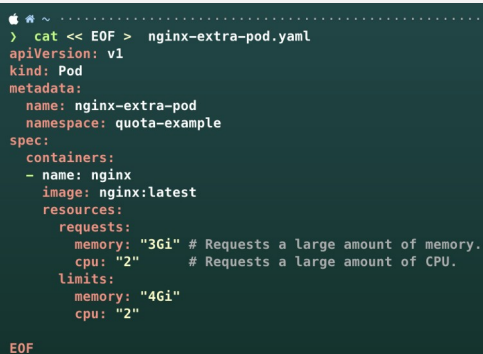
> kubectl describe pods -l app=nginx -n quota-example
Name:          nginx-replicaset-48stw
Namespace:     quota-example
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.3
Start Time:    Fri, 15 Nov 2024 14:28:21 +0530
Labels:        app=nginx
Annotations:    <none>
Status:        Running
IP:            10.1.0.23
IPs:           10.1.0.23
Controlled By: ReplicaSet/nginx-replicaset
Containers:
  nginx:
    Container ID:  docker://a47f94cd31a6ebf71b762779350a41f9cc9040f4e909f3bb2b70e4cc810525d3
    Image:         nginx:latest
    Image ID:      docker-pullable://nginx@sha256:bc5eac5eafc581aeda3008b4b1f07ebba230de2f27d47767129a6a905c84f470
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Fri, 15 Nov 2024 14:28:50 +0530
    Ready:         True
    Restart Count:  0
    Limits:
      cpu:    200m
      memory: 200Mi
    Requests:
      cpu:    100m
      memory: 100Mi
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-x8b94 (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized         True
  Ready               True
  ContainersReady      True
  PodScheduled        True
Volumes:
```

Attempt to Exceed the Resource Quota

Try creating additional resources to see if they are rejected when exceeding the quota. For example, create more Pods or increase the CPU/memory requests to exceed the quota limits.

Create a YAML file named ***nginx-extra-pod.yaml*** with the following content:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-extra-pod
  namespace: quota-example
spec:
  containers:
  - name: nginx
    image: nginx:latest
  resources:
    requests:
      memory: "3Gi" # Requests a large amount of memory.
      cpu: "2"      # Requests a large amount of CPU.
    limits:
      memory: "4Gi"
      cpu: "2"
```



```
Apple ~ % cat << EOF > nginx-extra-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: nginx-extra-pod
  namespace: quota-example
spec:
  containers:
  - name: nginx
    image: nginx:latest
  resources:
    requests:
      memory: "3Gi" # Requests a large amount of memory.
      cpu: "2"      # Requests a large amount of CPU.
    limits:
      memory: "4Gi"
      cpu: "2"
EOF
```


Apply this YAML to create the Pod:

kubectl apply -f nginx-extra-pod.yaml

```
❯ kubectl apply -f nginx-extra-pod.yaml
Error from server (Forbidden): error when creating "nginx-extra-pod.yaml": pods "nginx-extra-pod" is forbidden: exceeded quota: example-quota, requested: requests.cpu=2, used: requests.cpu=500m, limited: requests.cpu=2
```

This should fail due to exceeding the Resource Quota. Check the events to see the failure reason:

kubectl get events -n quota-example

```
❯ kubectl get events -n quota-example
```

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
3m49s	Normal	Scheduled	pod/nginx-replicaset-48stw	Successfully assigned quota-example/nginx-replicaset-48stw to docker-desktop
3m49s	Normal	Pulling	pod/nginx-replicaset-48stw	Pulling image "nginx:latest"
3m20s	Normal	Pulled	pod/nginx-replicaset-48stw	Successfully pulled image "nginx:latest" in 9.989s (29.02s including waiting). Image size: 196880043 bytes.
3m20s	Normal	Created	pod/nginx-replicaset-48stw	Created container nginx
3m20s	Normal	Started	pod/nginx-replicaset-48stw	Started container nginx
3m49s	Normal	Scheduled	pod/nginx-replicaset-52g9d	Successfully assigned quota-example/nginx-replicaset-52g9d to docker-desktop
3m49s	Normal	Pulling	pod/nginx-replicaset-52g9d	Pulling image "nginx:latest"
3m35s	Normal	Pulled	pod/nginx-replicaset-52g9d	Successfully pulled image "nginx:latest" in 5.002s (13.932s including waiting). Image size: 196880043 bytes.
3m35s	Normal	Created	pod/nginx-replicaset-52g9d	Created container nginx
3m35s	Normal	Started	pod/nginx-replicaset-52g9d	Started container nginx
3m40s	Normal	Scheduled	pod/nginx-replicaset-bl956	Successfully assigned quota-example/nginx-replicaset-bl956 to docker-desktop
3m40s	Normal	Pulling	pod/nginx-replicaset-bl956	Pulling image "nginx:latest"
3m30s	Normal	Pulled	pod/nginx-replicaset-bl956	Successfully pulled image "nginx:latest" in 5.262s (19.194s including waiting). Image size: 196880043 bytes.
3m30s	Normal	Created	pod/nginx-replicaset-bl956	Created container nginx
3m30s	Normal	Started	pod/nginx-replicaset-bl956	Started container nginx
3m49s	Normal	Scheduled	pod/nginx-replicaset-lrnjn	Successfully assigned quota-example/nginx-replicaset-lrnjn to docker-desktop
3m49s	Normal	Pulling	pod/nginx-replicaset-lrnjn	Pulling image "nginx:latest"
3m40s	Normal	Pulled	pod/nginx-replicaset-lrnjn	Successfully pulled image "nginx:latest" in 5.285s (8.93s including waiting). Image size: 196880043 bytes.
3m40s	Normal	Created	pod/nginx-replicaset-lrnjn	Created container nginx
3m40s	Normal	Started	pod/nginx-replicaset-lrnjn	Started container nginx
3m49s	Normal	Scheduled	pod/nginx-replicaset-p2w8d	Successfully assigned quota-example/nginx-replicaset-p2w8d to docker-desktop
3m49s	Normal	Pulling	pod/nginx-replicaset-p2w8d	Pulling image "nginx:latest"
3m45s	Normal	Pulled	pod/nginx-replicaset-p2w8d	Successfully pulled image "nginx:latest" in 3.644s (3.644s including waiting). Image size: 196880043 bytes.
3m45s	Normal	Created	pod/nginx-replicaset-p2w8d	Created container nginx
3m45s	Normal	Started	pod/nginx-replicaset-p2w8d	Started container nginx
3m50s	Normal	SuccessfulCreate	replicaset/nginx-replicaset	Created pod: nginx-replicaset-52g9d
3m49s	Normal	SuccessfulCreate	replicaset/nginx-replicaset	Created pod: nginx-replicaset-bl956
3m49s	Normal	SuccessfulCreate	replicaset/nginx-replicaset	Created pod: nginx-replicaset-lrnjn
3m49s	Normal	SuccessfulCreate	replicaset/nginx-replicaset	Created pod: nginx-replicaset-p2w8d
3m49s	Normal	SuccessfulCreate	replicaset/nginx-replicaset	Created pod: nginx-replicaset-48stw

Look for error messages indicating that the Pod creation was denied due to resource constraints.

Step 6: Clean Up Resources

To delete the resources you created:

```
kubectl delete -f nginx-replicaset-quota.yaml
```

```
kubectl delete -f nginx-extra-pod.yaml
```

```
kubectl delete -f resource-quota.yaml
```

```
kubectl delete namespace quota-example
```

```
🍏 🏠 ~ .....  
> kubectl delete -f nginx-replicaset-quota.yaml  
kubectl delete -f nginx-extra-pod.yaml  
kubectl delete -f resource-quota.yaml  
kubectl delete namespace quota-example  
replicaset.apps "nginx-replicaset" deleted  
Error from server (NotFound): error when deleting "nginx-extra-pod.yaml": pods "nginx-extra-pod" not found  
resourcequota "example-quota" deleted  
namespace "quota-example" deleted
```