

EXPERIMENT 8

AIM: Creating and Managing a ReplicaSet in Kubernetes

Objective:

A ReplicaSet in Kubernetes ensures a specified number of Pod replicas are running at any given time. This exercise will guide you through creating a ReplicaSet to maintain the desired state of your application.

- Understand the syntax and structure of a Kubernetes ReplicaSet definition file (YAML).
- Learn how to create and manage a ReplicaSet to ensure application availability.
- Understand how a ReplicaSet helps in scaling applications and maintaining desired states.

Prerequisites

- Kubernetes Cluster: Have a running Kubernetes cluster (locally using Minikube or kind, or a cloud-based service).
- kubectl: Install and configure kubectl to interact with your Kubernetes cluster.
- Basic Knowledge of YAML: Familiarity with YAML format will be helpful for understanding Kubernetes resource definitions.

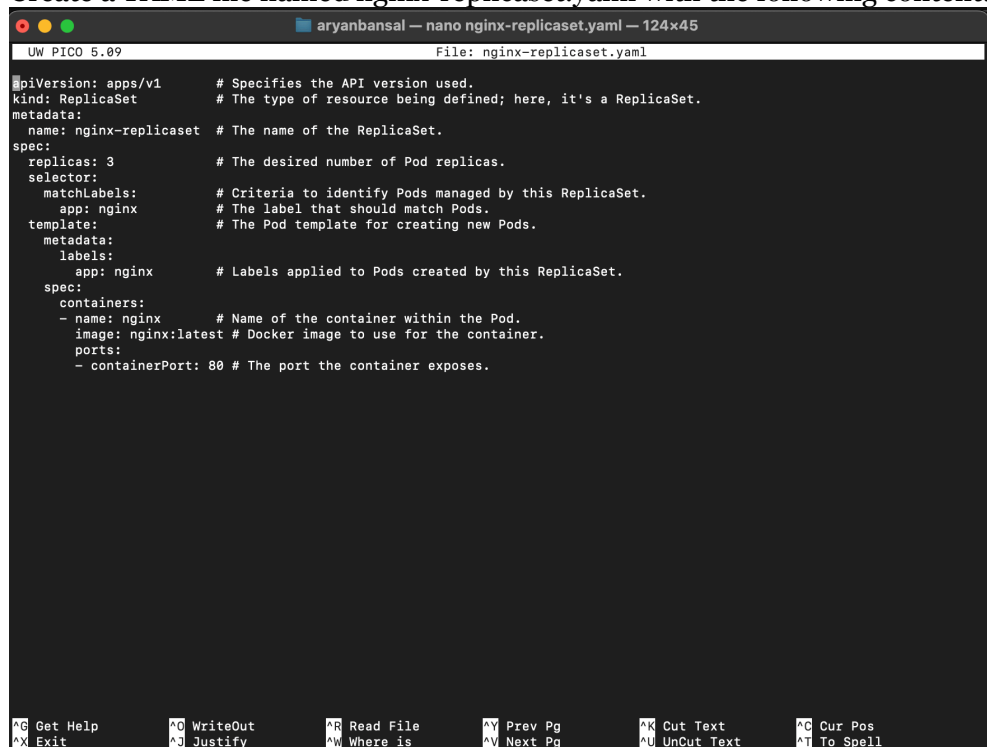
Step-by-Step Guide

Step 1: Understanding ReplicaSet

A ReplicaSet ensures a specified number of Pod replicas are running at any given time. If a Pod crashes or is deleted, the ReplicaSet creates a new one to meet the defined number of replicas. This helps maintain application availability and ensures that your application can handle increased load by distributing traffic among multiple Pods.

Step 2: Create a ReplicaSet

We'll define a ReplicaSet to maintain three replicas of a simple Nginx web server Pod. Create a YAML file named `nginx-replicaset.yaml` with the following content:



```
apiVersion: apps/v1      # Specifies the API version used.
kind: ReplicaSet         # The type of resource being defined; here, it's a ReplicaSet.
metadata:
  name: nginx-replicaset # The name of the ReplicaSet.
spec:
  replicas: 3            # The desired number of Pod replicas.
  selector:
    matchLabels:         # Criteria to identify Pods managed by this ReplicaSet.
      app: nginx         # The label that should match Pods.
  template:              # The Pod template for creating new Pods.
    metadata:
      labels:
        app: nginx       # Labels applied to Pods created by this ReplicaSet.
    spec:
      containers:
        - name: nginx     # Name of the container within the Pod.
          image: nginx:latest # Docker image to use for the container.
          ports:
            - containerPort: 80 # The port the container exposes.
```

Explanation:

- **apiVersion:** Defines the API version (apps/v1) used for the ReplicaSet resource.
- **kind:** Specifies that this resource is a ReplicaSet.
- **metadata:** Contains metadata about the ReplicaSet, including name.
- **name:** The unique name for the ReplicaSet.
- **spec:** Provides the specification for the ReplicaSet.
- **replicas:** Defines the desired number of Pod replicas.
- **selector:** Criteria for selecting Pods managed by this ReplicaSet.
- **matchLabels:** Labels that Pods must have to be managed by this ReplicaSet.
- **template:** Defines the Pod template used for creating new Pods.
- **metadata:** Contains metadata for the Pods, including labels.
- **labels:** Labels applied to Pods created by this ReplicaSet.
- **ospec:** Specification for the Pods.
- **containers:** Lists the containers that will run in the Pod.
- **name:** The unique name of the container within the Pod.
- **image:** The Docker image used for the container.
- **ports:** Ports exposed by the container.

Step 3: Apply the YAML to Create the ReplicaSet

Use the **kubectl apply** command to create the ReplicaSet based on the YAML file.

kubectl apply -f nginx-replicaset.yaml

```
Done! Kubectl is now configured to use minikube cluster and default namespace by default
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f nginx-replicaset.yaml
replicaset.apps/nginx-replicaset created
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Verify the ReplicaSet is running and maintaining the desired number of replicas:

kubectl get replicaset

```
replicaset.apps/nginx-replicaset created
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get replicaset
NAME                DESIRED   CURRENT   READY   AGE
nginx-replicaset    3         3         0       41m
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

This command lists all ReplicaSets in the current namespace.

To check the Pods created by the ReplicaSet:

kubectl get pods -l app=nginx

This command lists all Pods with the label **app=nginx**.

```
nginx-replicaset    3         3         0       41m
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get pods -l app=nginx
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-d817c             1/1     Running   0           42m
nginx-replicaset-ghrxf             0/1     ImagePullBackOff  0           42m
nginx-replicaset-pfmzp             1/1     Running   0           42m
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Step 4: Managing the ReplicaSet

1. Scaling the ReplicaSet

You can scale the number of replicas managed by the ReplicaSet using the **kubectl scale** command.

kubectl scale --replicas=5 replicaset/nginx-replicaset

This command scales the ReplicaSet to maintain 5 replicas. Verify the scaling operation:

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl scale --replicas=5 replicaset/nginx-replicaset
replicaset.apps/nginx-replicaset scaled
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

```
kubectl get pods -l app=nginx
```

You should see that the number of Pods has increased to 5.

```
replicaset.apps/nginx-replicaset scaled
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get pods -l app=nginx
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-d8l7c              1/1     Running   0           44m
nginx-replicaset-dwrnx              1/1     Running   0           48s
nginx-replicaset-ghrxf              1/1     Running   0           44m
nginx-replicaset-kvpg7              1/1     Running   0           48s
nginx-replicaset-pfmzp              1/1     Running   0           44m
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

2. Updating the ReplicaSet

If you need to update the Pod template (e.g., to use a different Docker image version), modify the YAML file and apply it again. For instance, change the image to a specific version of Nginx:

Apply the changes:

```
kubectl apply -f nginx-replicaset.yaml
```

```
UW PICO 5.09                                     File: nginx-replicaset.yaml

spec:
  template:
    spec:
      containers:
      - name: nginx
        image: nginx:1.19.3 # Change to a specific version

^G Get Help      ^O WriteOut      ^R Read File     ^Y Prev Pg
^X Exit          ^J Justify       ^W Where is      ^V Next Pg

(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl apply -f nginx-replicaset.yaml
replicaset.apps/nginx-replicaset configured
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Check the status to ensure the Pods are updated:

```
kubectl get pods -l app=nginx
```

```
(base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl get pods -l app=nginx
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-d8l7c              1/1     Running   2 (2m39s ago)  57m
nginx-replicaset-ghrxf              1/1     Running   2 (2m39s ago)  57m
nginx-replicaset-kvpg7              1/1     Running   2 (2m39s ago)  13m
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```

Note: Updating a ReplicaSet doesn't automatically replace existing Pods with new ones. In practice, you often create a new ReplicaSet or Deployment for updates.

3. Deleting the ReplicaSet

To clean up the ReplicaSet and its Pods, use the kubectl delete command:

`kubectl delete -f nginx-replicaset.yaml`

```
((base) aryanbansal@Aryans-MacBook-Air-10 ~ % kubectl delete -f nginx-replicaset.yaml
replicaset.apps "nginx-replicaset" deleted
(base) aryanbansal@Aryans-MacBook-Air-10 ~ %
```