

Lab Exercise 10- Implementing Resource Quota in Kubernetes

Objective:

In Kubernetes, Resource Quotas are used to control the resource consumption of namespaces. They help in managing and enforcing limits on the usage of resources like CPU, memory, and the number of objects (e.g., Pods, Services) within a namespace. This exercise will guide you through creating and managing Resource Quotas to limit the resources used by applications in a specific namespace.

Step 1: Understand Resource Quotas

Resource Quotas allow you to:

- Limit the amount of CPU and memory a namespace can use.
- Control the number of certain types of resources (e.g., Pods, Services, PersistentVolumeClaims) in a namespace.
- Prevent a namespace from consuming more resources than allocated, ensuring fair usage across multiple teams or applications.

Step 2: Create a Namespace

First, create a namespace where you will apply the Resource Quota. This helps in isolating and controlling resource usage within that specific namespace.

Create a YAML file named ***quota-namespace.yaml*** with the following content:

```
apiVersion: v1
kind: Namespace
metadata:
  name: quota-example # The name of the namespace.
```

Apply the YAML to create the namespace:

```
kubectl apply -f quota-namespace.yaml
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl apply -f namespace.yaml
namespace/quota-example created
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> |
```

Verify that the namespace is created:

```
kubectl get namespaces
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl get namespaces
NAME                STATUS    AGE
default             Active   38h
kube-node-lease     Active   38h
kube-public         Active   38h
kube-system         Active   38h
quota-example       Active   23s
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> |
```

You should see quota-example listed in the output.

Step 3: Define a Resource Quota

Next, create a Resource Quota YAML file named ***resource-quota.yaml*** with the following content:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: example-quota # The name of the Resource Quota.
  namespace: quota-example # The namespace to which the Resource Quota will apply.
spec:
  hard:
    # The hard limits imposed by this Resource Quota.
    requests.cpu: "2" # The total CPU resource requests allowed in the namespace (2 cores).
    requests.memory: "4Gi" # The total memory resource requests allowed in the namespace (4 GiB).
    limits.cpu: "4" # The total CPU resource limits allowed in the namespace (4 cores).
    limits.memory: "8Gi" # The total memory resource limits allowed in the namespace (8 GiB).
    pods: "10" # The total number of Pods allowed in the namespace.
    persistentvolumeclaims: "5" # The total number of PersistentVolumeClaims allowed in the namespace.
    configmaps: "10" # The total number of ConfigMaps allowed in the namespace.
    services: "5" # The total number of Services allowed in the namespace.
```

Step 4: Apply the Resource Quota

Apply the Resource Quota YAML to the namespace:

```
kubectl apply -f resource-quota.yaml
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> kubectl apply -f resourceQuota.yaml
resourcequota/example-quota created
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> |
```

Verify that the Resource Quota is applied:

```
kubectl get resourcequota -n quota-example
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> kubectl get resourcequota -n quota-example
NAME          AGE   REQUEST
LIMIT
example-quota 47s   configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4
Gi, services: 0/5 limits.cpu: 0/4, limits.memory: 0/8Gi
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> |
```

To see the details of the applied Resource Quota:

```
kubectl describe resourcequota example-quota -n quota-example
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> kubectl describe resourcequota example-quota -n quota-example
Name:          example-quota
Namespace:     quota-example
Resource       Used  Hard
-----
configmaps     1    10
limits.cpu     0     4
limits.memory  0    8Gi
persistentvolumeclaims 0     5
pods           0    10
requests.cpu   0     2
requests.memory 0    4Gi
services       0     5
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docke_lab\lab10> |
```

Step 5: Test the Resource Quota

Let's create some resources in the quota-example namespace to see how the Resource Quota affects them.

Deploy a ReplicaSet with Resource Requests and Limits

Create a YAML file named ***nginx-replicaset-quota.yaml*** with the following content:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  namespace: quota-example
spec:
  replicas: 5          # Desired number of Pod replicas.
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
          resources:    # Define resource requests and limits.
            requests:
              memory: "100Mi"
              cpu: "100m"
            limits:
              memory: "200Mi"
              cpu: "200m"
```

Explanation:

This ReplicaSet requests a total of 500m CPU and 500Mi memory across 5 replicas. It also limits each replica to use a maximum of 200m CPU and 200Mi memory.

Apply this YAML to create the ReplicaSet:

```
kubectl apply -f nginx-replicaset-quota.yaml
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl apply -f replica.yaml
replicaset.apps/nginx-replicaset created
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> |
```

Check the status of the Pods and ensure they are created within the constraints of the Resource Quota:

```
kubectl get pods -n quota-example
```

```
replicaset.apps/nginx-replicaset created
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl get pods -n quota-example
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-48cd5              1/1     Running   0           33s
nginx-replicaset-5g9xw              1/1     Running   0           33s
nginx-replicaset-bxfnf              1/1     Running   0           33s
nginx-replicaset-lrvf2              1/1     Running   0           33s
nginx-replicaset-pvf25              1/1     Running   0           33s
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> |
```

To describe the Pods and see their resource allocations:

```
kubectl describe pods -l app=nginx -n quota-example
```

```

PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl describe resourcequota example-quota -n quota-example
Name:          example-quota
Namespace:     quota-example
Resource       Used  Hard
-----
configmaps    1    10
limits.cpu    0     4
limits.memory 0    8Gi
persistentvolumeclaims 0     5
pods          0    10
requests.cpu  0     2
requests.memory 0    4Gi
services      0     5
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl apply -f replica.yaml
replicaset.apps/nginx-replicaset created
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl get pods -n quota-example
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-48cd5             1/1     Running   0           33s
nginx-replicaset-5g9xw             1/1     Running   0           33s
nginx-replicaset-bxfnf             1/1     Running   0           33s
nginx-replicaset-lrvf2             1/1     Running   0           33s
nginx-replicaset-pvf25             1/1     Running   0           33s
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl describe pods -l app=nginx -n example-quota
No resources found in example-quota namespace.

```

Attempt to Exceed the Resource Quota

Try creating additional resources to see if they are rejected when exceeding the quota. For example, create more Pods or increase the CPU/memory requests to exceed the quota limits.

Create a YAML file named ***nginx-extra-pod.yaml*** with the following content:

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx-extra-pod
  namespace: quota-example
spec:
  containers:
  - name: nginx
    image: nginx:latest
    resources:
      requests:
        memory: "3Gi" # Requests a large amount of memory.
        cpu: "2"      # Requests a large amount of CPU.
      limits:
        memory: "4Gi"
        cpu: "2"

```

Apply this YAML to create the Pod:

```
kubectl apply -f nginx-extra-pod.yaml
```

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl apply -f extra-pod.yaml
Error from server (Forbidden): error when creating "extra-pod.yaml": pods "nginx-extra-pod" is forbidden: exceeded quota
: example-quota, requested: requests.cpu=2, used: requests.cpu=500m, limited: requests.cpu=2
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> |
```

This should fail due to exceeding the Resource Quota. Check the events to see the failure reason:

```
kubectl get events -n quota-example
```

Look for error messages indicating that the Pod creation was denied due to resource constraints.

Step 6: Clean Up Resources

To delete the resources you created:

```
PS C:\Users\KHUSHI JAIN\OneDrive\Desktop\Docker_lab\lab10> kubectl delete namespace quota-example
namespace "quota-example" deleted
```

```
kubectl delete -f nginx-replicaset-quota.yaml
```

```
kubectl delete -f nginx-extra-pod.yaml
```

```
kubectl delete -f resource-quota.yaml
```

```
kubectl delete namespace quota-example
```